

# Home Credit



on a établi un modèle pour prédire dans quelle mesure chaque client est capable de rembourser un prêt.

## Concept

Dans ce problème, les données sont déséquilibrées. Nous ne pouvons donc pas utiliser la précision comme mesure d'erreur. On va utiliser :

- Log loss, F1-score et AUC
- La courbe Roc (visualiser les performances des classificateur binaire)
- La Matrice de confusion (Obtenir un aperçu des prévisions)

## Les 7 sources de données (Datasets)

**Application\_train/test:** les principales base de données pour chaque demande de crédit et chaque prêt identifié par (SK\_ID\_CURR). TARGET indiquant 1 si le prêt n'a pas été remboursé ou 0 si le prêt a été remboursé. Dans le cas de notre modèle on utilise seulement les données d'entraînement (**Application\_train**)

**bureau:** cet ensemble de données comprends les crédits précédent auprès d'autres institutions financières (client's previous credits).

**bureau\_balance:** se compose de données mensuelles sur les crédits précédents, Chaque ligne correspond à un mois d'un crédit précédent et un seul dépend de chaque mois de la durée du crédit .

**previous\_application :** Les données des demandes précédentes de prêts au crédit immobilier des clients. Chaque application précédente a une ligne et est identifiée par la fonction (SK\_ID\_PREV).

**POS\_CASH\_BALANCE :** se compose de données mensuelles sur les points de vente précédents ou les prêts de trésorerie que les clients ont obtenus avec le crédit immobilier.

**credit\_card\_balance:** Les données mensuelles sur les anciennes cartes de crédit que les clients avaient avec Home Credit.

**installments\_payment:** l'historique de paiement des prêts anciens chez Home Credit.

les étapes des datasets :

1. Vérifier combien de colonnes et de lignes
2. La jointure des Datasets :
  - Previous Application data et Application Bureau data
  - POS\_CASH\_balance data et application\_bureau\_prev\_data
  - Installments Payments data et application\_bureau\_prev\_dat
  - a
3. Répartition des datasets en train, test et valid
4. Normalisée notre dataset
5. La sélection des features

## Modèles de Machine Learning

1. LightGBM : est un cadre de gradient rapide, distribué et haute performance basé sur un algorithme d'arbre de décision, utilisé pour le classement, la classification et de nombreuses autres tâches d'apprentissage automatique.

2. Logistic Regression :

3. Random Forest Classifier

# Table des matières

1. Importer les librairies
2. Lire le Dataset
3. Visualization (EDA)
4. Un model de LightGBM
5. Un model de régression logistique
6. Un model de Random Forest Classifier

## 1. Importer les librairies

```
Entrée [49]: #Importer les librairies
import pandas as pd
import sklearn
import numpy as np
import matplotlib.pyplot as plt
import os
import warnings
import seaborn as sns
from sklearn.preprocessing import OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler
from sklearn.svm import LinearSVC
from sklearn.metrics import roc_auc_score
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_auc_score
from sklearn.calibration import CalibratedClassifierCV
from sklearn.metrics import confusion_matrix
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.linear_model import SGDClassifier
import plotly.offline as py
import plotly.graph_objs as go
from plotly.offline import init_notebook_mode, iplot
from sklearn.model_selection import train_test_split
init_notebook_mode(connected=True)
import cufflinks as cf
cf.go_offline()
import pickle
import gc
import lightgbm as lgb
from lightgbm import LGBMClassifier
warnings.filterwarnings('ignore')
import matplotlib.pyplot as plt
%matplotlib inline
#Installation de cufflinks
#pip install cufflinks
#pip install lightgbm
```

## 2. Lire le Dataset

```
Entrée [2]: #Lire le dataset Train
application= pd.read_csv("/Users/macbook/Downloads/application_train.csv")
print('done!!!')
print('The shape of data:',application.shape)
print('First 5 rows of data:')
application.head()
```

done!!!  
The shape of data: (307511, 122)  
First 5 rows of data:

Out[2]:

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY
0	100002	1	Cash loans	M	N	
1	100003	0	Cash loans	F	N	
2	100004	0	Revolving loans	M	Y	
3	100006	0	Cash loans	F	N	
4	100007	0	Cash loans	M	N	

5 rows × 122 columns

```
Entrée [3]: #Calculer les valeurs manquantes avec des pourcentages
count = application.isnull().sum().sort_values(ascending=False)
percentage = ((application.isnull().sum()/len(application)*100)).sort
missing_application = pd.concat([count, percentage], axis=1, keys=[ 'C
print('Count and percentage of missing values for top 20 columns:')
missing_application.head(20)
```

Count and percentage of missing values for top 20 columns:

Out[3]:


	Count	Percentage
<b>COMMONAREA_MEDI</b>	214865	69.872297
<b>COMMONAREA_AVG</b>	214865	69.872297
<b>COMMONAREA_MODE</b>	214865	69.872297
<b>NONLIVINGAPARTMENTS_MODE</b>	213514	69.432963
<b>NONLIVINGAPARTMENTS_MEDI</b>	213514	69.432963
<b>NONLIVINGAPARTMENTS_AVG</b>	213514	69.432963
<b>FONDKAPREMONT_MODE</b>	210295	68.386172
<b>LIVINGAPARTMENTS_MEDI</b>	210199	68.354953
<b>LIVINGAPARTMENTS_MODE</b>	210199	68.354953
<b>LIVINGAPARTMENTS_AVG</b>	210199	68.354953
<b>FLOORSMIN_MEDI</b>	208642	67.848630
<b>FLOORSMIN_MODE</b>	208642	67.848630
<b>FLOORSMIN_AVG</b>	208642	67.848630
<b>YEARS_BUILD_MEDI</b>	204488	66.497784
<b>YEARS_BUILD_AVG</b>	204488	66.497784
<b>YEARS_BUILD_MODE</b>	204488	66.497784
<b>OWN_CAR_AGE</b>	202929	65.990810
<b>LANDAREA_MODE</b>	182590	59.376738
<b>LANDAREA_AVG</b>	182590	59.376738
<b>LANDAREA_MEDI</b>	182590	59.376738

```
Entrée [4]: columns_without_id = [col for col in application.columns if col != 'SK_
#Vérifier les duplicates dans notre base
application[application.duplicated(subset = columns_without_id, keep=
print('The no of duplicates in the data:', application[application.dup
.shape[0])]
```

The no of duplicates in the data: 0

### 3. Visualization (EDA)


```
Entrée [5]: #Le type de credit
cf.set_config_file(theme='polar')
contract_val = application['NAME_CONTRACT_TYPE'].value_counts()
contract_df = pd.DataFrame({'labels': contract_val.index,
                            'values': contract_val.values
                           })
contract_df.iplot(kind='pie', labels='labels', values='values', title='
```



Entrée [6]: #La distribution des revenus

```
application[application['AMT_INCOME_TOTAL'] < 2000000]['AMT_INCOME_TOTAL']
xTitle = 'Total Income', yTitle ='Count of applicants',
title='Distribution of AMT_INCOME_TOTAL')
```

Distribution of AMT\_INCOME\_TOTAL



Entrée [7]: #Calculer le nombre des erreurs pour Target et Days Employed

```
error = application[application['DAYS_EMPLOYED'] == 365243]
print('The no of errors are :', len(error))
(error['TARGET'].value_counts()/len(error))*100
```

The no of errors are : 55374

Out[7]:

Days Employed	Percentage
0	94.600354
1	5.399646

Name: TARGET, dtype: float64

Entrée [8]: # Crée une colonne

```
application['DAYS_EMPLOYED_ERROR'] = application["DAYS_EMPLOYED"] ==
# Replace the error values with nan
application['DAYS_EMPLOYED'].replace({365243: np.nan}, inplace = True)
```

```
Entrée [9]: cf.set_config_file(theme='pearl')
(application['DAYS_EMPLOYED']/(-365)).iplot(kind='histogram', xTitle
                                         yTitle='Count of applicants in %',
                                         title='Years before the application the person started c
```

Years before the application the person started current employ



```
Entrée [10]: application[application['DAYS_EMPLOYED']>(-365*2)]['TARGET'].value_c
```

```
Out[10]: 0      0.887924
          1      0.112076
Name: TARGET, dtype: float64
```

## Preparation de notre data :

```
Entrée [11]: Représentation de Total income quand c'est supérieur à Credit
application['INCOME_GT_CREDIT_FLAG'] = application['AMT_INCOME_TOTAL'] >
    Colonne de pourcentage de Credit Income Percent
application['CREDIT_INCOME_PERCENT'] = application['AMT_CREDIT'] / application['AMT_INCOME_TOTAL'] * 100 >
    Column to represent Annuity Income percent
application['ANNUITY_INCOME_PERCENT'] = application['AMT_ANNUITY'] / application['AMT_INCOME_TOTAL'] * 100 >
    Colonne qui représente les Credit Term
application['CREDIT_TERM'] = application['AMT_CREDIT'] / application['AMT_ANNUITY'] >
    Colonne qui représente les Days Employed pourcentage
application['DAYS_EMPLOYED_PERCENT'] = application['DAYS_EMPLOYED'] / application['AMT_ANNUITY'] * 100 >
    Shape of Application data
cint('The shape of application data:', application.shape)
```

The shape of application data: (307511, 128)

```
Entrée [12]: print('Reading the data....', end='')
bureau = pd.read_csv("/Users/macbook/Downloads/bureau.csv")
print('done!!!')
print('The shape of data:', bureau.shape)
print('First 5 rows of data:')
bureau.head()
```

Reading the data....done!!!
The shape of data: (1716428, 17)
First 5 rows of data:

Out[12]:

	SK_ID_CURR	SK_ID_BUREAU	CREDIT_ACTIVE	CREDIT_CURRENCY	DAYS_CREDIT	CRE
0	215354	5714462	Closed	currency 1	-497	
1	215354	5714463	Active	currency 1	-208	
2	215354	5714464	Active	currency 1	-203	
3	215354	5714465	Active	currency 1	-203	
4	215354	5714466	Active	currency 1	-629	

```
Entrée [13]: # Combiner les variables numériques (features)
grp = bureau.drop(['SK_ID_BUREAU'], axis = 1).groupby(by=['SK_ID_CURR'])
grp.columns = ['BUREAU_'+column if column !='SK_ID_CURR' else column
application_bureau = application.merge(grp, on='SK_ID_CURR', how='left')
application_bureau.update(application_bureau[grp.columns].fillna(0))

# Combiner les variable catégoriques
bureau_categorical = pd.get_dummies(bureau.select_dtypes('object'))
bureau_categorical['SK_ID_CURR'] = bureau['SK_ID_CURR']
grp = bureau_categorical.groupby(by = ['SK_ID_CURR']).mean().reset_index()
grp.columns = ['BUREAU_'+column if column !='SK_ID_CURR' else column
application_bureau = application_bureau.merge(grp, on='SK_ID_CURR', how='left')
application_bureau.update(application_bureau[grp.columns].fillna(0))

# Shape de application et bureau data combiner
print('The shape application and bureau data combined:', application_bureau.shape)
```

The shape application and bureau data combined: (307511, 163)

```
Entrée [14]: # Numéro des anciens crédits par client
grp = bureau.groupby(by = ['SK_ID_CURR'])['SK_ID_BUREAU'].count().reset_index()
application_bureau = application_bureau.merge(grp, on='SK_ID_CURR', how='left')
application_bureau['BUREAU_LOAN_COUNT'] = application_bureau['BUREAU_SK_ID_CURR']

# Numéro de type des anciens crédits par client
grp = bureau[['SK_ID_CURR', 'CREDIT_TYPE']].groupby(by = ['SK_ID_CURR'])
application_bureau = application_bureau.merge(grp, on='SK_ID_CURR', how='left')
application_bureau['BUREAU_LOAN_TYPES'] = application_bureau['BUREAU_SK_ID_CURR']

# Ratio dette/ crédit
bureau['AMT_CREDIT_SUM'] = bureau['AMT_CREDIT_SUM'].fillna(0)
bureau['AMT_CREDIT_SUM_DEBT'] = bureau['AMT_CREDIT_SUM_DEBT'].fillna(0)
grp1 = bureau[['SK_ID_CURR', 'AMT_CREDIT_SUM']].groupby(by=['SK_ID_CURR'])
grp2 = bureau[['SK_ID_CURR', 'AMT_CREDIT_SUM_DEBT']].groupby(by=['SK_ID_CURR'])
grp1['DEBT_CREDIT_RATIO'] = grp2['TOTAL_CREDIT_SUM_DEBT']/grp1['TOTAL_CREDIT_SUM']
del grp1['TOTAL_CREDIT_SUM']
application_bureau = application_bureau.merge(grp1, on='SK_ID_CURR', how='left')
application_bureau['DEBT_CREDIT_RATIO'] = application_bureau['DEBT_CREDIT_RATIO'].replace([None], 0)
application_bureau['DEBT_CREDIT_RATIO'] = pd.to_numeric(application_bureau['DEBT_CREDIT_RATIO'])

# Le ratio d'endettement
bureau['AMT_CREDIT_SUM_OVERDUE'] = bureau['AMT_CREDIT_SUM_OVERDUE'].fillna(0)
bureau['AMT_CREDIT_SUM_DEBT'] = bureau['AMT_CREDIT_SUM_DEBT'].fillna(0)
grp1 = bureau[['SK_ID_CURR', 'AMT_CREDIT_SUM_OVERDUE']].groupby(by=['SK_ID_CURR'])
grp2 = bureau[['SK_ID_CURR', 'AMT_CREDIT_SUM_DEBT']].groupby(by=['SK_ID_CURR'])
grp1['OVERDUE_DEBT_RATIO'] = grp1['TOTAL_CUSTOMER_OVERDUE']/grp2['TOTAL_CUSTOMER_DEBT']
del grp1['TOTAL_CUSTOMER_OVERDUE']
application_bureau = application_bureau.merge(grp1, on='SK_ID_CURR', how='left')
application_bureau['OVERDUE_DEBT_RATIO'] = application_bureau['OVERDUE_DEBT_RATIO'].replace([None], 0)
application_bureau['OVERDUE_DEBT_RATIO'] = pd.to_numeric(application_bureau['OVERDUE_DEBT_RATIO'])
```

```
Entrée [15]: print('Reading the data....', end='')
previous_applicaton = pd.read_csv("/Users/macbook/Downloads/previous
print('The shape of data:', previous_applicaton.shape)
print('First 5 rows of data:')
previous_applicaton.head()
```

Reading the data....The shape of data: (1670214, 37)  
First 5 rows of data:

Out[15]:

	SK_ID_PREV	SK_ID_CURR	NAME_CONTRACT_TYPE	AMT_ANNUITY	AMT_APPLICATION	A
0	2030495	271877	Consumer loans	1730.430	17145.0	
1	2802425	108129	Cash loans	25188.615	607500.0	
2	2523466	122040	Cash loans	15060.735	112500.0	
3	2819243	176158	Cash loans	47041.335	450000.0	
4	1784265	202054	Cash loans	31924.395	337500.0	

5 rows × 37 columns

```
Entrée [16]: vious applications per customer
applicaton[['SK_ID_CURR', 'SK_ID_PREV']].groupby(by=['SK_ID_CURR'])['SK_ID_PREV'].count().reset_index()
eau_prev = application_bureau.merge(grp, on=['SK_ID_CURR'], how='left')
eau_prev['PREV_APP_COUNT'] = application_bureau_prev['PREV_APP_COUNT']
variables numériques (features)
applicaton.drop(['SK_ID_PREV'], axis=1).groupby(by=['SK_ID_CURR']).mean()
['PREV_'+column if column != 'SK_ID_CURR' else column for column in grp.columns]
rev_columns
eau_prev = application_bureau_prev.merge(grp, on=['SK_ID_CURR'], how='left')
eau_prev.update(application_bureau_prev[grp.columns].fillna(0))
variables catégoriques
l = pd.get_dummies(previous_applicaton.select_dtypes('object'))
l['SK_ID_CURR'] = previous_applicaton['SK_ID_CURR']
l.head()
gorical.groupby('SK_ID_CURR').mean().reset_index()
['PREV_'+column if column != 'SK_ID_CURR' else column for column in grp.columns]
eau_prev = application_bureau_prev.merge(grp, on=['SK_ID_CURR'], how='left')
eau_prev.update(application_bureau_prev[grp.columns].fillna(0))
```

```
Entrée [17]: print('Reading the data....', end='')
pos_cash = pd.read_csv("/Users/macbook/Downloads/POS_CASH_balance.csv")
print('The shape of data:', pos_cash.shape)
print('First 5 rows of data:')
pos_cash.head()
```

Reading the data....The shape of data: (10001358, 8)  
First 5 rows of data:

Out[17]:

	SK_ID_PREV	SK_ID_CURR	MONTHS_BALANCE	CNT_INSTALMENT	CNT_INSTALMENT_FU1
0	1803195	182943	-31	48.0	
1	1715348	367990	-33	36.0	
2	1784872	397406	-32	12.0	
3	1903291	269225	-35	48.0	
4	2341044	334279	-35	36.0	

```
Entrée [18]: #inser les variable numériques (features)
pos_cash.drop('SK_ID_PREV', axis=1).groupby(by=['SK_ID_CURR']).mean()
columns = ['POS_'+column if column != 'SK_ID_CURR' else column for column in columns]
columns = prev_columns
application_bureau_prev = application_bureau_prev.merge(grp, on=['SK_ID_CURR'])
application_bureau_prev.update(application_bureau_prev[grp.columns].fillna(0))
#inser les variable catégoriques
sh_categorical = pd.get_dummies(pos_cash.select_dtypes('object'))
sh_categorical['SK_ID_CURR'] = pos_cash['SK_ID_CURR']
pos_cash_categorical.groupby('SK_ID_CURR').mean().reset_index()
columns = ['POS_'+column if column != 'SK_ID_CURR' else column for column in columns]
application_bureau_prev = application_bureau_prev.merge(grp, on=['SK_ID_CURR'])
application_bureau_prev.update(application_bureau_prev[grp.columns].fillna(0))
```

```
Entrée [19]: print('Reading the data....', end='')
insta_payments = pd.read_csv("/Users/macbook/Downloads/installments_
print('The shape of data:', insta_payments.shape)
print('First 5 rows of data:')
insta_payments.head()
```

Reading the data....The shape of data: (13605401, 8)  
First 5 rows of data:

Out[19]:

	SK_ID_PREV	SK_ID_CURR	NUM_INSTALMENT_VERSION	NUM_INSTALMENT_NUMBER	DA
0	1054186	161674		1.0	6
1	1330831	151639		0.0	34
2	2085231	193053		2.0	1
3	2452527	199697		1.0	3
4	2714724	167756		1.0	2

```
Entrée [20]: Combiner les variables numériques et pas catégoriques
cp = insta_payments.drop('SK_ID_PREV', axis=1).groupby(by=['SK_ID_CU
prev_columns = ['INSTA_' + column if column != 'SK_ID_CURR' else column
cp.columns = prev_columns
application_bureau_prev = application_bureau_prev.merge(grp, on=['SK_
application_bureau_prev.update(application_bureau_prev[grp.columns].fi
```

```
Entrée [21]: print('Reading the data....', end='')
credit_card = pd.read_csv("/Users/macbook/Downloads/credit_card_bala
print('The shape of data:', credit_card.shape)
print('First 5 rows of data:')
credit_card.head()
```

Reading the data....The shape of data: (3840312, 23)  
First 5 rows of data:

Out[21]:

	SK_ID_PREV	SK_ID_CURR	MONTHS_BALANCE	AMT_BALANCE	AMT_CREDIT_LIMIT_ACTL
0	2562384	378907		-6	56.970
1	2582071	363914		-1	63975.555
2	1740877	371185		-7	31815.225
3	1389973	337855		-4	236572.110
4	1891521	126868		-1	453919.455

5 rows × 23 columns

```
Entrée [22]: # Combiner les variables numériques (features)
grp = credit_card.drop('SK_ID_PREV', axis=1).groupby(by=['SK_ID_CURR'])
prev_columns = ['CREDIT_' + column if column != 'SK_ID_CURR' else column
for column in grp.columns]
grp.columns = prev_columns
application_bureau_prev = application_bureau_prev.merge(grp, on=['SK_ID_CURR'])
application_bureau_prev.update(application_bureau_prev[grp.columns])
# Combiner les variables catégoriques
credit_categorical = pd.get_dummies(credit_card.select_dtypes('object'))
credit_categorical['SK_ID_CURR'] = credit_card['SK_ID_CURR']
grp = credit_categorical.groupby('SK_ID_CURR').mean().reset_index()
grp.columns = ['CREDIT_' + column if column != 'SK_ID_CURR' else column
for column in grp.columns]
application_bureau_prev = application_bureau_prev.merge(grp, on=['SK_ID_CURR'])
application_bureau_prev.update(application_bureau_prev[grp.columns])
```

```
Entrée [23]: X_train, X_val, X_test, y_train, y_val, y_test = train_test_split(application_bureau_prev,
                                                               X_temp, y_temp, stratify=y_temp)
print('Shape of X_train:', X_train.shape)
print('Shape of X_val:', X_val.shape)
print('Shape of X_test:', X_test.shape)
```

```
Shape of X_train: (215257, 375)
Shape of X_val: (46127, 375)
Shape of X_test: (46127, 375)
```

```
Entrée [24]: # Séparation des colonnes numériques et catégoriques
types = np.array([dt for dt in X_train.dtypes])
all_columns = X_train.columns.values
is_num = types != 'object'
num_cols = all_columns[is_num]
cat_cols = all_columns[~is_num]
```

```
Entrée [25]: # Caractérisation des données numériques
imputer_num = SimpleImputer(strategy='median')
X_train_num = imputer_num.fit_transform(X_train[num_cols])
X_val_num = imputer_num.transform(X_val[num_cols])
X_test_num = imputer_num.transform(X_test[num_cols])
scaler_num = StandardScaler()
X_train_num1 = scaler_num.fit_transform(X_train_num)
X_val_num1 = scaler_num.transform(X_val_num)
X_test_num1 = scaler_num.transform(X_test_num)
X_train_num_final = pd.DataFrame(X_train_num1, columns=num_cols)
X_val_num_final = pd.DataFrame(X_val_num1, columns=num_cols)
X_test_num_final = pd.DataFrame(X_test_num1, columns=num_cols)

# Caractérisation des données catégoriques
imputer_cat = SimpleImputer(strategy='constant', fill_value='MISSING')
X_train_cat = imputer_cat.fit_transform(X_train[cat_cols])
X_val_cat = imputer_cat.transform(X_val[cat_cols])
X_test_cat = imputer_cat.transform(X_test[cat_cols])
X_train_cat1 = pd.DataFrame(X_train_cat, columns=cat_cols)
X_val_cat1 = pd.DataFrame(X_val_cat, columns=cat_cols)
X_test_cat1 = pd.DataFrame(X_test_cat, columns=cat_cols)
ohe = OneHotEncoder(sparse=False, handle_unknown='ignore')
X_train_cat2 = ohe.fit_transform(X_train_cat1)
X_val_cat2 = ohe.transform(X_val_cat1)
X_test_cat2 = ohe.transform(X_test_cat1)
cat_cols_ohe = list(ohe.get_feature_names(input_features=cat_cols))
X_train_cat_final = pd.DataFrame(X_train_cat2, columns = cat_cols_ohe)
X_val_cat_final = pd.DataFrame(X_val_cat2, columns = cat_cols_ohe)
X_test_cat_final = pd.DataFrame(X_test_cat2, columns = cat_cols_ohe)
# Données finales
X_train_final = pd.concat([X_train_num_final,X_train_cat_final], axis=1)
X_val_final = pd.concat([X_val_num_final,X_val_cat_final], axis = 1)
X_test_final = pd.concat([X_test_num_final,X_test_cat_final], axis = 1)
print(X_train_final.shape)
print(X_val_final.shape)
print(X_test_final.shape)
```

```
(215257, 505)
(46127, 505)
(46127, 505)
```

## LGBMClassifier

```
Entrée [26]: model_sk = lgb.LGBMClassifier(boosting_type='gbdt', max_depth=7, lea
                     class_weight='balanced', subsample=0.9, colsample_b
train_features, valid_features, train_y, valid_y = train_test_split(
model_sk.fit(train_features, train_y, early_stopping_rounds=100, eva
```

```
Training until validation scores don't improve for 100 rounds
[200]  valid_0's auc: 0.75423  valid_0's binary_logloss: 0.592408
[400]  valid_0's auc: 0.768815  valid_0's binary_logloss: 0.566125
[600]  valid_0's auc: 0.774772  valid_0's binary_logloss: 0.551609
[800]  valid_0's auc: 0.777189  valid_0's binary_logloss: 0.541956
[1000] valid_0's auc: 0.778678  valid_0's binary_logloss: 0.534552
[1200] valid_0's auc: 0.77957   valid_0's binary_logloss: 0.52803
[1400] valid_0's auc: 0.779734  valid_0's binary_logloss: 0.522452
Early stopping, best iteration is:
[1332] valid_0's auc: 0.779798  valid_0's binary_logloss: 0.524251
```

```
Out[26]: LGBMClassifier(boosting_type='gbdt', class_weight='balanced',
                      colsample_bytree=0.8, importance_type='split',
                      learning_rate=0.01, max_depth=7, min_child_samples=20
                      ,
                      min_child_weight=0.001, min_split_gain=0.0, n_estimators=2000,
                      n_jobs=-1, num_leaves=31, objective=None, random_state=None,
                      reg_alpha=0.0, reg_lambda=0.0, silent=True, subsample=0.9,
                      subsample_for_bin=200000, subsample_freq=0)
```

```
Entrée [32]: feature_imp = pd.DataFrame(sorted(zip(model_sk.feature_importances_, X
features_df = feature_imp.sort_values(by="Value", ascending=False)
selected_features = list(features_df[features_df['Value']>=50]['Feature
Enregistrement des fonctionnalités sélectionnées dans un fichier pic
ith open('select_features.txt','wb') as fp:
    pickle.dump(selected_features, fp)
cint('The no. of features selected:',len(selected_features))
```

```
ERROR! Session/line number was not unique in database. History loggi
ng moved to new session 1392
The no. of features selected: 179
```

## Regression Logistique

```
Entrée [33]: def plot_confusion_matrix(test_y, predicted_y):
    # Confusion matrix
    C = confusion_matrix(test_y, predicted_y)

    # Recall matrix
    A = (((C.T)/(C.sum(axis=1))).T)

    # Precision matrix
    B = (C/(C.sum(axis=0)))

    plt.figure(figsize=(20,4))

    labels = ['Re-paid(0)', 'Not Re-paid(1)']
    cmap=sns.light_palette("purple")
    plt.subplot(1,3,1)
    sns.heatmap(C, annot=True, cmap=cmap, fmt="d", xticklabels = labels,
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title('Confusion matrix')

    plt.subplot(1,3,2)
    sns.heatmap(A, annot=True, cmap=cmap, xticklabels = labels, ytic
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title('Recall matrix')

    plt.subplot(1,3,3)
    sns.heatmap(B, annot=True, cmap=cmap, xticklabels = labels, ytic
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title('Precision matrix')

    plt.show()
def cv_plot(alpha, cv_auc):

    fig, ax = plt.subplots()
    ax.plot(np.log10(alpha), cv_auc,c='g')
    for i, txt in enumerate(np.round(cv_auc,3)):
        ax.annotate((alpha[i],str(txt)), (np.log10(alpha[i]),cv_auc[
    plt.grid()
    plt.xticks(np.log10(alpha))
    plt.title("Cross Validation Error for each alpha")
    plt.xlabel("Alpha i's")
    plt.ylabel("Error measure")
    plt.show()
```

```
Entrée [42]: alpha = np.logspace(-4,4,9)
cv_auc_score = []
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l1', class_weight = 'balanc
```

```
Entrée [35]: clf.fit(X_train_final[selected_features], y_train)
```

```
Out[35]: SGDClassifier(alpha=10000.0, average=False, class_weight='balanced',
                       early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
                       l1_ratio=0.15, learning_rate='optimal', loss='log', max_iter=1000,
                       n_iter_no_change=5, n_jobs=None, penalty='l1', power_t=0.5,
                       random_state=28, shuffle=True, tol=0.001, validation_fraction=0.1,
                       verbose=0, warm_start=False)
```

```
Entrée [36]: sig_clf = CalibratedClassifierCV(clf, method='sigmoid')
sig_clf.fit(X_train_final[selected_features], y_train)
```


```
Out[36]: CalibratedClassifierCV(base_estimator=SGDClassifier(alpha=10000.0,
                                                               average=False,
                                                               class_weight='balanced',
                                                               early_stopping=False,
                                                               epsilon=0.1, eta0=0.0,
                                                               fit_intercept=True,
                                                               l1_ratio=0.15,
                                                               learning_rate='optimal',
                                                               loss='log', max_iter=1000,
                                                               n_iter_no_change=5,
                                                               n_jobs=None, penalty='l1',
                                                               power_t=0.5,
                                                               random_state=28,
                                                               shuffle=True, tol=0.001,
                                                               validation_fraction=0.1,
                                                               verbose=0,
                                                               warm_start=False),
                                 cv=None, method='sigmoid')
```

```
Entrée [37]: y_pred_prob = []
sig_clf.predict_proba(X_val_final[selected_features])[:,1]
```


```
Out[37]: array([0.08072677, 0.08072677, 0.08072677, ..., 0.08072677, 0.08072677,
   0.08072677])
```

```
Entrée [40]: best_alpha = alpha[np.argmax(cv_auc_score)]
logreg = SGDClassifier(alpha = best_alpha, class_weight = 'balanced')
logreg.fit(X_train_final[selected_features], y_train)
logreg_sig_clf = CalibratedClassifierCV(logreg, method='sigmoid')
logreg_sig_clf.fit(X_train_final[selected_features], y_train)
y_pred_prob = logreg_sig_clf.predict_proba(X_train_final[selected_features])
print('For best alpha {0}, The Train AUC score is {1}'.format(best_alpha))
y_pred_prob = logreg_sig_clf.predict_proba(X_val_final[selected_features])
print('For best alpha {0}, The Cross validated AUC score is {1}'.format(best_alpha))
y_pred_prob = logreg_sig_clf.predict_proba(X_test_final[selected_features])
print('For best alpha {0}, The Test AUC score is {1}'.format(best_alpha))
y_pred = logreg.predict(X_test_final[selected_features])
print('The test AUC score is :', roc_auc_score(y_test,y_pred_prob))
print('The percentage of misclassified points {:.05.2f}% :'.format((1 - roc_auc_score(y_test,y_pred)) * 100))
plot_confusion_matrix(y_test, y_pred)
```

For best alpha 0.0001, The Train AUC score is 0.7650162581393501  
 For best alpha 0.0001, The Cross validated AUC score is 0.7569473024589161  
 For best alpha 0.0001, The Test AUC score is 0.7604168247220617  
 The test AUC score is : 0.7604168247220617  
 The percentage of misclassified points 32.31% :



```
Entrée [41]: from sklearn.metrics import roc_curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
auc = roc_auc_score(y_test,y_pred_prob)
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, marker='.')
plt.plot([0, 1], [0, 1], linestyle='--')
plt.title('ROC curve', fontsize = 20)
plt.xlabel('FPR', fontsize=15)
plt.ylabel('TPR', fontsize=15)
plt.grid()
plt.legend([ "AUC=% .3f" %auc])
plt.show()
```



## RandomForestClassifier

```
Entrée [ ]: lpha = [200,500,1000,2000]
ax_depth = [7, 10]
v_auc_score = []
or i in alpha:
    for j in max_depth:
        clf = RandomForestClassifier(n_estimators=i, criterion='gini',
                                      random_state=42, n_jobs=-1)
        clf.fit(X_train_final[selected_features], y_train)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(X_train_final[selected_features], y_train)
        y_pred_prob = sig_clf.predict_proba(X_val_final[selected_features])
        cv_auc_score.append(roc_auc_score(y_val,y_pred_prob))
    print('For n_estimators {0}, max_depth {1} cross validation AUC
          format(i,j,roc_auc_score(y_val,y_pred_prob)))
```

```
For n_estimators 200, max_depth 7 cross validation AUC score 0.74594
69826033478
For n_estimators 200, max_depth 10 cross validation AUC score 0.7509
055165092412
For n_estimators 500, max_depth 7 cross validation AUC score 0.74624
43061744536
For n_estimators 500, max_depth 10 cross validation AUC score 0.7511
577697532852
For n_estimators 1000, max_depth 7 cross validation AUC score 0.7462
187597785892
For n_estimators 1000, max_depth 10 cross validation AUC score 0.751
1696626961294
For n_estimators 2000, max_depth 7 cross validation AUC score 0.7462
969758260167
```